

Seminar Algorithms for Big Data

Fast Random Integer Generation in an Interval

Based on a paper of the same title by Daniel Lemire

Lukas Geis

Supervised by Dr. Manuel Penschuck

29th February 2024 · Algorithm Engineering (Prof. Dr. Ulrich Meyer)

What is our goal?



What is our goal?

We want to *efficiently* draw a *uniform* random integer in an interval.



What is our goal?

We want to *efficiently* draw a *uniform* random integer in an interval.

Where do we need this?



What is our goal?

We want to *efficiently* draw a *uniform* random integer in an interval.

Where do we need this?

- Shuffling



TBD

What is our goal?

We want to *efficiently* draw a *uniform* random integer in an interval.

Where do we need this?

- Shuffling
- Complex Graph Generators



TBD

TBD

What is our goal?

We want to *efficiently* draw a *uniform* random integer in an interval.

Where do we need this?

- Shuffling
- Complex Graph Generators
- Sampling

TBD

TBD

TBD

Table of Contents

1 Preliminaries

- Formal Definition
- Operations
- The Naive Approach

2 Unbiased Algorithms

- The OpenBSD Algorithm
- The Java Algorithm
- The Bitmask Algorithm

3 Lemire's Algorithm

- Multiply-And-Shift
- The Algorithm

4 Conclusion

1 Preliminaries





Formal Definition

Setting:



Formal Definition

Setting:

- **Input:** upper bound of interval $n \in \mathbb{N}$



Formal Definition

Setting:

- **Input:** upper bound of interval $n \in \mathbb{N}$
- **Output:** uniform random integer in interval $[0, n)$



Formal Definition

Setting:

- **Input:** upper bound of interval $n \in \mathbb{N}$
- **Output:** uniform random integer in interval $[0, n)$

But what if we want a random integer in $[a, b)$ for $a, b \in \mathbb{N}$, $0 < a < b$ instead?

Formal Definition

Setting:

- **Input:** upper bound of interval $n \in \mathbb{N}$
- **Output:** uniform random integer in interval $[0, n)$

But what if we want a random integer in $[a, b)$ for $a, b \in \mathbb{N}$, $0 < a < b$ instead?

We can map this to our setting by subtracting a !

Formal Definition

Setting:

- **Input:** upper bound of interval $n \in \mathbb{N}$
- **Output:** uniform random integer in interval $[0, n)$

But what if we want a random integer in $[a, b)$ for $a, b \in \mathbb{N}$, $0 < a < b$ instead?

We can map this to our setting by subtracting a !

- Set $n = b - a$ and draw a uniform random integer $x \in [0, n)$

Formal Definition

Setting:

- **Input:** upper bound of interval $n \in \mathbb{N}$
- **Output:** uniform random integer in interval $[0, n)$

But what if we want a random integer in $[a, b)$ for $a, b \in \mathbb{N}$, $0 < a < b$ instead?

We can map this to our setting by subtracting a !

- Set $n = b - a$ and draw a uniform random integer $x \in [0, n)$
- Return $x + a$



Definition (Common Operations)



Definition (Common Operations)

■ Integer-Division: $x \div y \quad := \lfloor x/y \rfloor$



Definition (Common Operations)

- Integer-Division: $x \div y \quad := \lfloor x/y \rfloor$
- Remainder-Operation: $x \bmod y \quad := x - (x \div y)y$

Definition (Common Operations)

- Integer-Division: $x \div y \quad := \lfloor x/y \rfloor$
- Remainder-Operation: $x \bmod y \quad := x - (x \div y)y$
- Bit-RIGHTSHIFT: $x \gg W \quad := x \div 2^W$

Definition (Common Operations)

- Integer-Division: $x \div y := \lfloor x/y \rfloor$
- Remainder-Operation: $x \bmod y := x - (x \div y)y$
- Bit-RIGHTSHIFT: $x \gg W := x \div 2^W$
- Bit-LEFTSHIFT: $x \ll W := x \cdot 2^W$

Definition (Common Operations)

- Integer-Division: $x \div y := \lfloor x/y \rfloor$
- Remainder-Operation: $x \bmod y := x - (x \div y)y$
- Bit-RIGHTSHIFT: $x \gg W := x \div 2^W$
- Bit-LEFTSHIFT: $x \ll W := x \cdot 2^W$
- Bitwise-AND: $x \& y$

Definition (Common Operations)

- Integer-Division: $x \div y \quad := \lfloor x/y \rfloor$
- Remainder-Operation: $x \bmod y \quad := x - (x \div y)y$
- Bit-RIGHTSHIFT: $x \gg W \quad := x \div 2^W$
- Bit-LEFTSHIFT: $x \ll W \quad := x \cdot 2^W$
- Bitwise-AND: $x \& y \rightarrow x \bmod 2^W \quad := x \& (2^W - 1)$

Definition (Common Operations)

- Integer-Division: $x \div y := \lfloor x/y \rfloor$
- Remainder-Operation: $x \bmod y := x - (x \div y)y$
- Bit-RIGHTSHIFT: $x \gg W := x \div 2^W$
- Bit-LEFTSHIFT: $x \ll W := x \cdot 2^W$
- Bitwise-AND: $x \& y \rightarrow x \bmod 2^W := x \& (2^W - 1)$

Definition (Power Remainder)

For $W, n \in \mathbb{N}$, we write \mathcal{R}_n^W for $2^W \bmod n$.

The Naive Approach



The Naive Approach

How do we get random numbers?



The Naive Approach

How do we get random numbers?

- Generated by Pseudo-Random-Number-Generators (PRNGs)



The Naive Approach

How do we get random numbers?

- Generated by Pseudo-Random-Number-Generators (PRNGs)
- Generated as W -bit words, i.e. unsigned integers in $[0, 2^W)$ (typically $W \in \{32, 64\}$)



The Naive Approach

How do we get random numbers?

- Generated by Pseudo-Random-Number-Generators (PRNGs)
- Generated as W -bit words, i.e. unsigned integers in $[0, 2^W)$ (typically $W \in \{32, 64\}$)

`rand() mod n`



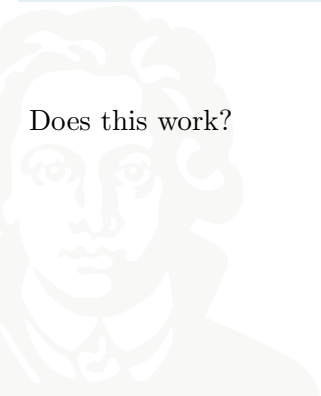
The Naive Approach

How do we get random numbers?

- Generated by Pseudo-Random-Number-Generators (PRNGs)
- Generated as W -bit words, i.e. unsigned integers in $[0, 2^W)$ (typically $W \in \{32, 64\}$)

`rand() mod n`

Does this work?



The Naive Approach

How do we get random numbers?

- Generated by Pseudo-Random-Number-Generators (PRNGs)
- Generated as W -bit words, i.e. unsigned integers in $[0, 2^W)$ (typically $W \in \{32, 64\}$)

`rand() mod n`

Does this work?

- Yes, the generated number is in $[0, n)$.

The Naive Approach

How do we get random numbers?

- Generated by Pseudo-Random-Number-Generators (PRNGs)
- Generated as W -bit words, i.e. unsigned integers in $[0, 2^W)$ (typically $W \in \{32, 64\}$)

`rand() mod n`

Does this work?

- Yes, the generated number is in $[0, n)$.

Is this efficient?

The Naive Approach

How do we get random numbers?

- Generated by Pseudo-Random-Number-Generators (PRNGs)
- Generated as W -bit words, i.e. unsigned integers in $[0, 2^W)$ (typically $W \in \{32, 64\}$)

$\text{rand()} \bmod n$

Does this work?

- Yes, the generated number is in $[0, n)$.

Is this efficient?

- No, we require one expensive integer division operation.

The Naive Approach

How do we get random numbers?

- Generated by Pseudo-Random-Number-Generators (PRNGs)
- Generated as W -bit words, i.e. unsigned integers in $[0, 2^W)$ (typically $W \in \{32, 64\}$)

$\text{rand()} \bmod n$

Does this work?

- Yes, the generated number is in $[0, n)$.

Is this efficient?

- No, we require one expensive integer division operation.

Is the generated number uniform in $[0, n)$?

The Naive Approach



The Naive Approach

In general, applying $x \bmod n$ to $[0, 2^W)$ yields



The Naive Approach

In general, applying $x \bmod n$ to $[0, 2^W)$ yields

$$\underbrace{\underbrace{0, 1, \dots, n-1}_{n \text{ values}}, \underbrace{0, 1, \dots, n-1}_{n \text{ values}}, \dots, \underbrace{0, 1, \dots, n-1}_{n \text{ values}}}_{(2^W \div n) \cdot n \text{ values}}, \underbrace{0, 1, \dots, \mathcal{R}_n^W - 1}_{\mathcal{R}_n^W \text{ values}}$$

2^W values

The Naive Approach

In general, applying $x \bmod n$ to $[0, 2^W)$ yields

$$\underbrace{\underbrace{0, 1, \dots, n-1}_{n \text{ values}}, \underbrace{0, 1, \dots, n-1}_{n \text{ values}}, \dots, \underbrace{0, 1, \dots, n-1}_{n \text{ values}}}_{(2^W \div n) \cdot n \text{ values}}, \underbrace{0, 1, \dots, \mathcal{R}_n^W - 1}_{\mathcal{R}_n^W \text{ values}}$$

2^W values

We have a **leftover** interval that introduces bias.

The Naive Approach

In general, applying $x \bmod n$ to $[0, 2^W)$ yields

$$\underbrace{\underbrace{0, 1, \dots, n-1}_{n \text{ values}}, \underbrace{0, 1, \dots, n-1}_{n \text{ values}}, \dots, \underbrace{0, 1, \dots, n-1}_{n \text{ values}}}_{(2^W \div n) \cdot n \text{ values}}, \underbrace{0, 1, \dots, \mathcal{R}_n^W - 1}_{\mathcal{R}_n^W \text{ values}}$$

2^W values

We have a **leftover** interval that introduces bias.

Every approach that maps every integer in $[0, 2^W)$ to a single number in $[0, n)$

The Naive Approach

In general, applying $x \bmod n$ to $[0, 2^W)$ yields

$$\underbrace{\underbrace{0, 1, \dots, n-1}_{n \text{ values}}, \underbrace{0, 1, \dots, n-1}_{n \text{ values}}, \dots, \underbrace{0, 1, \dots, n-1}_{n \text{ values}}}_{(2^W \div n) \cdot n \text{ values}}, \underbrace{0, 1, \dots, \mathcal{R}_n^W - 1}_{\mathcal{R}_n^W \text{ values}}$$

2^W values

We have a **leftover** interval that introduces bias.

Every approach that maps every integer in $[0, 2^W)$ to a single number in $[0, n)$ does **not** generate uniform random integers in one step

The Naive Approach

In general, applying $x \bmod n$ to $[0, 2^W)$ yields

$$\underbrace{\underbrace{0, 1, \dots, n-1}_{n \text{ values}}, \underbrace{0, 1, \dots, n-1}_{n \text{ values}}, \dots, \underbrace{0, 1, \dots, n-1}_{n \text{ values}}}_{(2^W \div n) \cdot n \text{ values}}, \underbrace{0, 1, \dots, \mathcal{R}_n^W - 1}_{\mathcal{R}_n^W \text{ values}}$$

2^W values

We have a **leftover** interval that introduces bias.

Every approach that maps every integer in $[0, 2^W)$ to a single number in $[0, n)$ does **not** generate uniform random integers in one step whenever n does not divide 2^W .

The Naive Approach

In general, applying $x \bmod n$ to $[0, 2^W)$ yields

$$\underbrace{\overbrace{0, 1, \dots, n-1}^{n \text{ values}}, \overbrace{0, 1, \dots, n-1}^{n \text{ values}}, \dots, \overbrace{0, 1, \dots, n-1}^{n \text{ values}}}_{(2^W \div n) \cdot n \text{ values}}, \underbrace{0, 1, \dots, \mathcal{R}_n^W - 1}_{\mathcal{R}_n^W \text{ values}}$$

We have a **leftover** interval that introduces bias.

Every approach that maps every integer in $[0, 2^W)$ to a single number in $[0, n)$ does **not** generate uniform random integers in one step whenever n does not divide 2^W .

Idea: Use **rejection sampling** to achieve uniformity!

2

Unbiased Algorithms



The OpenBSD Algorithm



The OpenBSD Algorithm

- We shift the **rejection interval** to the left:



The OpenBSD Algorithm

- We shift the **rejection interval** to the left:

$$\underbrace{\underbrace{0, 1, \dots, \mathcal{R}_n^W - 1}_{\mathcal{R}_n^W \text{ values}}, \underbrace{\mathcal{R}_n^W, \dots, n-1, 0, \dots, \mathcal{R}_n^W - 1}_{n \text{ values}}, \dots, \underbrace{\mathcal{R}_n^W, \dots, n-1, 0, \dots, \mathcal{R}_n^W - 1}_{n \text{ values}}}_{2^W \text{ values}}$$

$(2^W \div n) \cdot n \text{ values}$

The OpenBSD Algorithm

- We shift the **rejection interval** to the left:

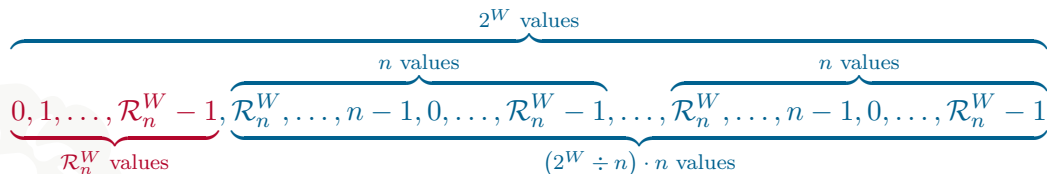
$$\overbrace{
 \underbrace{0, 1, \dots, \mathcal{R}_n^W - 1}_{\mathcal{R}_n^W \text{ values}},
 \underbrace{\mathcal{R}_n^W, \dots, n-1, 0, \dots, \mathcal{R}_n^W - 1}_{n \text{ values}},
 \underbrace{\mathcal{R}_n^W, \dots, n-1, 0, \dots, \mathcal{R}_n^W - 1}_{n \text{ values}}
 }^{2^W \text{ values}}$$

$(2^W \div n) \cdot n \text{ values}$

- Algorithm:

The OpenBSD Algorithm

- We shift the **rejection interval** to the left:

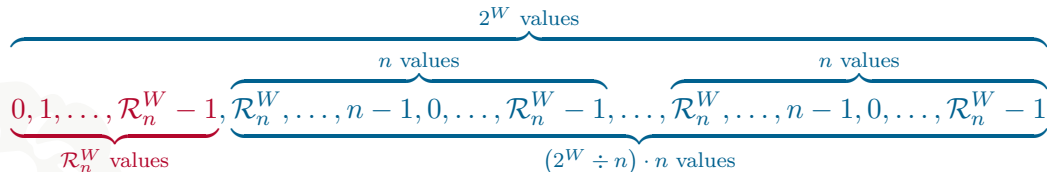


- Algorithm:

- Generate a uniform random number $x \in [0, 2^W)$ until $x \geq \mathcal{R}_n^W$

The OpenBSD Algorithm

- We shift the **rejection interval** to the left:



- Algorithm:

- Generate a uniform random number $x \in [0, 2^W)$ until $x \geq \mathcal{R}_n^W$
- Return $x \bmod n$

The OpenBSD Algorithm

Algorithm:

- Generate a uniform random number $x \in [0, 2^W)$ until $x \geq \mathcal{R}_n^W$
- Return $x \bmod n$



The OpenBSD Algorithm

Algorithm:

- Generate a uniform random number $x \in [0, 2^W)$ until $x \geq \mathcal{R}_n^W$
- Return $x \bmod n$

Efficiency

The OpenBSD Algorithm

Algorithm:

- Generate a uniform random number $x \in [0, 2^W)$ until $x \geq \mathcal{R}_n^W$
- Return $x \bmod n$

Efficiency

We require 2 integer division operations:

The OpenBSD Algorithm

Algorithm:

- Generate a uniform random number $x \in [0, 2^W)$ until $x \geq \mathcal{R}_n^W$
- Return $x \bmod n$

Efficiency

We require 2 integer division operations:

- one for computing \mathcal{R}_n^W

The OpenBSD Algorithm

Algorithm:

- Generate a uniform random number $x \in [0, 2^W)$ until $x \geq \mathcal{R}_n^W$
- Return $x \bmod n$

Efficiency

We require 2 integer division operations:

- one for computing \mathcal{R}_n^W
- and one for computing $x \bmod n$.

The OpenBSD Algorithm

Algorithm:

- Generate a uniform random number $x \in [0, 2^W)$ until $x \geq \mathcal{R}_n^W$
- Return $x \bmod n$

Efficiency

We require 2 integer division operations:

- one for computing \mathcal{R}_n^W
- and one for computing $x \bmod n$.

Can we do better?

The Java Algorithm



The Java Algorithm

- Consider $x - (x \bmod n)$ for $x \in [0, 2^W)$:



The Java Algorithm

- Consider $x - (x \bmod n)$ for $x \in [0, 2^W)$:

$$\begin{array}{c}
 \overbrace{\hspace{15em}}^{(2^W \div n) \cdot n \text{ values}} \\
 \underbrace{0, 1, \dots, n-1}_{n \text{ values}}, \underbrace{0, 1, \dots, n-1}_{n \text{ values}}, \dots, \underbrace{0, 1, \dots, n-1}_{n \text{ values}}, \underbrace{0, 1, \dots, \mathcal{R}_n^W - 1}_{\mathcal{R}_n^W \text{ values}}
 \end{array}$$

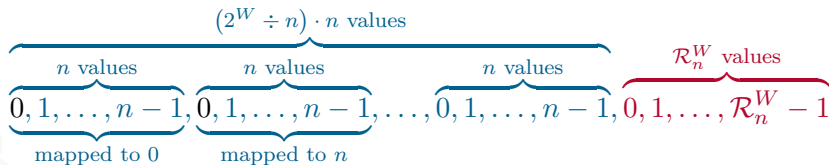
The Java Algorithm

- Consider $x - (x \bmod n)$ for $x \in [0, 2^W)$:

$$\begin{array}{c}
 \overbrace{\hspace{15em}}^{(2^W \div n) \cdot n \text{ values}} \\
 \underbrace{\overbrace{0, 1, \dots, n-1}^{n \text{ values}}, \overbrace{0, 1, \dots, n-1}^{n \text{ values}}, \dots, \overbrace{0, 1, \dots, n-1}^{n \text{ values}}}_{\text{mapped to 0}}, \overbrace{0, 1, \dots, \mathcal{R}_n^W - 1}^{\mathcal{R}_n^W \text{ values}}
 \end{array}$$

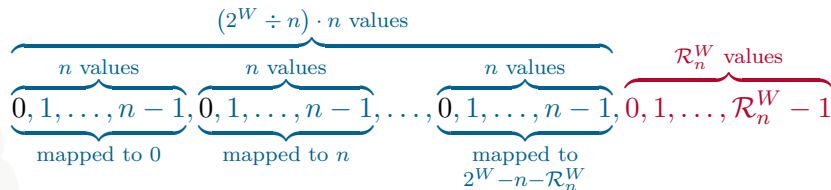
The Java Algorithm

- Consider $x - (x \bmod n)$ for $x \in [0, 2^W)$:



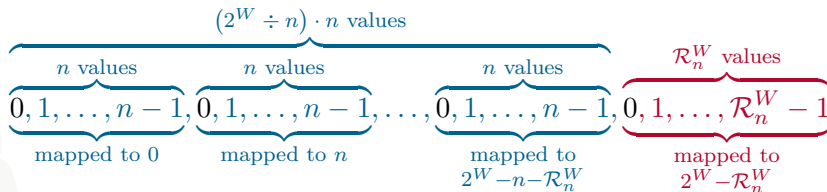
The Java Algorithm

- Consider $x - (x \bmod n)$ for $x \in [0, 2^W)$:



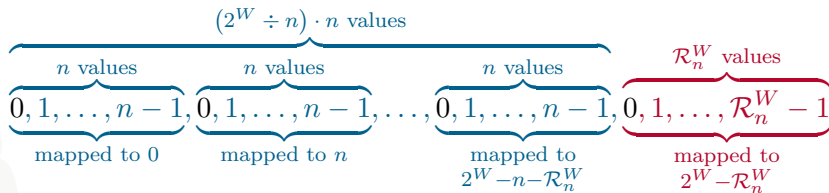
The Java Algorithm

- Consider $x - (x \bmod n)$ for $x \in [0, 2^W)$:



The Java Algorithm

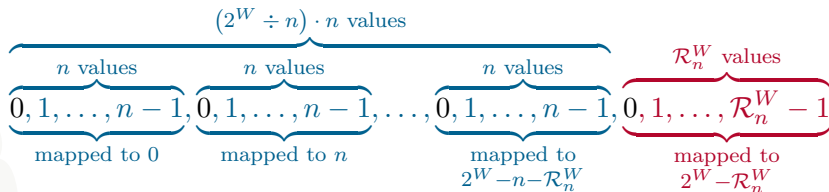
- Consider $x - (x \bmod n)$ for $x \in [0, 2^W)$:



- We map every number to the next-smallest multiple of n

The Java Algorithm

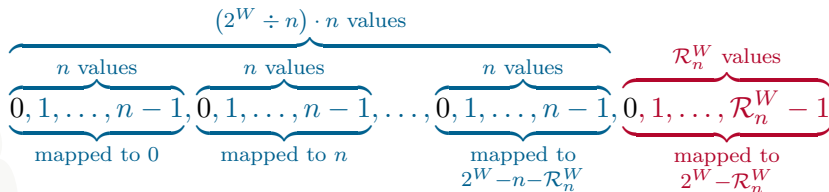
- Consider $x - (x \bmod n)$ for $x \in [0, 2^W)$:



- We map every number to the next-smallest multiple of n
- Only numbers in the **leftover** interval get mapped to $2^W - \mathcal{R}_n^W > 2^W - n$

The Java Algorithm

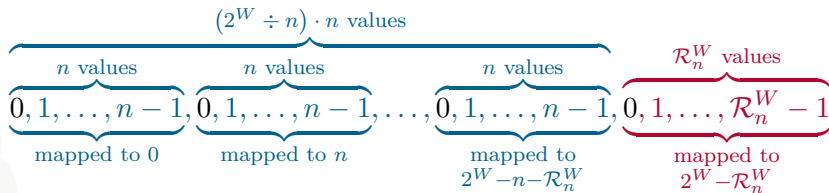
- Consider $x - (x \bmod n)$ for $x \in [0, 2^W)$:



- We map every number to the next-smallest multiple of n
- Only numbers in the **leftover** interval get mapped to $2^W - \mathcal{R}_n^W > 2^W - n$
- Algorithm:

The Java Algorithm

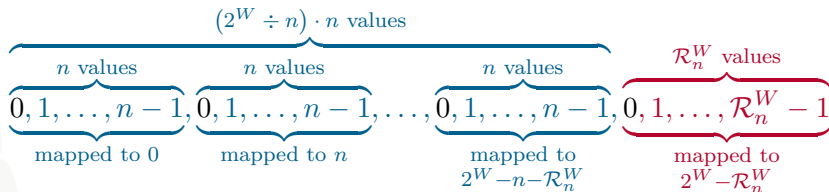
- Consider $x - (x \bmod n)$ for $x \in [0, 2^W)$:



- We map every number to the next-smallest multiple of n
- Only numbers in the **leftover** interval get mapped to $2^W - \mathcal{R}_n^W > 2^W - n$
- Algorithm:
 - (1) Draw $x \in [0, 2^W)$ and compute $r = x \bmod n$

The Java Algorithm

- Consider $x - (x \bmod n)$ for $x \in [0, 2^W)$:



- We map every number to the next-smallest multiple of n
- Only numbers in the **leftover** interval get mapped to $2^W - \mathcal{R}_n^W > 2^W - n$
- Algorithm:
 - Draw $x \in [0, 2^W)$ and compute $r = x \bmod n$
 - Return r if $x - r > 2^W - n$ else goto (1)

The Java Algorithm

Algorithm:

- (1) Draw $x \in [0, 2^W)$ and compute $r = x \bmod n$
- (2) Return r if $x - r > 2^W - n$ else goto (1)



The Java Algorithm

Algorithm:

- (1) Draw $x \in [0, 2^W)$ and compute $r = x \bmod n$
- (2) Return r if $x - r > 2^W - n$ else goto (1)

Efficiency

The Java Algorithm

Algorithm:

- (1) Draw $x \in [0, 2^W)$ and compute $r = x \bmod n$
- (2) Return r if $x - r > 2^W - n$ else goto (1)

Efficiency

- At least one integer division operation

The Java Algorithm

Algorithm:

- (1) Draw $x \in [0, 2^W)$ and compute $r = x \bmod n$
- (2) Return r if $x - r > 2^W - n$ else goto (1)

Efficiency

- At least one integer division operation
- Number of integer divisions operations equal to number of rounds

The Java Algorithm

Algorithm:

- (1) Draw $x \in [0, 2^W)$ and compute $r = x \bmod n$
- (2) Return r if $x - r > 2^W - n$ else goto (1)

Efficiency

- At least one integer division operation
- Number of integer divisions operations equal to number of rounds
- Return a number in round if $x < 2^W - \mathcal{R}_n^W$

The Java Algorithm

Algorithm:

- (1) Draw $x \in [0, 2^W)$ and compute $r = x \bmod n$
- (2) Return r if $x - r > 2^W - n$ else goto (1)

Efficiency

- At least one integer division operation
- Number of integer divisions operations equal to number of rounds
- Return a number in round if $x < 2^W - \mathcal{R}_n^W$
- Happens with probability $\frac{2^W - \mathcal{R}_n^W}{2^W} > \frac{1}{2}$

The Java Algorithm

Algorithm:

- (1) Draw $x \in [0, 2^W)$ and compute $r = x \bmod n$
- (2) Return r if $x - r > 2^W - n$ else goto (1)

Efficiency

- At least one integer division operation
- Number of integer divisions operations equal to number of rounds
- Return a number in round if $x < 2^W - \mathcal{R}_n^W$
- Happens with probability $\frac{2^W - \mathcal{R}_n^W}{2^W} > \frac{1}{2}$
- Expected number of integer division operations is $\frac{2^W}{2^W - \mathcal{R}_n^W} < 2$

The Bitmask Algorithm



The Bitmask Algorithm

- Consider the **binary** representation of n :



The Bitmask Algorithm

- Consider the binary representation of n :

$$\begin{array}{c}
 n \xrightarrow{\text{binary}} \underbrace{\overbrace{0, \dots, 0}^{2^{W-1}}}_{\text{only 0's}}, \underbrace{1}_{2^{\lfloor \log_2 n \rfloor}}, \underbrace{0/1, \dots, 0/1}_{\text{series of 0's and 1's}}^{\overbrace{2^1} \quad \overbrace{2^0}}
 \end{array}$$



The Bitmask Algorithm

- Consider the binary representation of n :

$$\begin{array}{c}
 n \xrightarrow{\text{binary}} \underbrace{\overbrace{0, \dots, 0}^{2^{W-1}}}_{\text{only 0's}}, \overbrace{1}^{2^{\lfloor \log_2 n \rfloor}}, \underbrace{0/1, \dots, \overbrace{0/1}^{2^1}, \overbrace{0/1}^{2^0}}_{\text{series of 0's and 1's}}
 \end{array}$$

- Every number $x \leq n$ only needs the last $\lfloor \log_2 n \rfloor + 1$ bits

The Bitmask Algorithm

- Consider the binary representation of n :

$$n \xrightarrow{\text{binary}} \underbrace{\overbrace{0, \dots, 0}^{2^{W-1}}}_{\text{only 0's}}, \overbrace{1}^{2^{\lfloor \log_2 n \rfloor}}, \underbrace{0/1, \dots, \overbrace{0/1}^{2^1}, \overbrace{0/1}^{2^0}}_{\text{series of 0's and 1's}}$$

- Every number $x \leq n$ only needs the last $\lfloor \log_2 n \rfloor + 1$ bits
- Get these bits with a bitwise-AND with

$$2^{\lfloor \log_2 n \rfloor + 1} - 1 \xrightarrow{\text{binary}} \underbrace{\overbrace{0, \dots, 0}^{2^{W-1}}}_{\text{only 0's}}, \underbrace{\overbrace{1}^{2^{\lfloor \log_2 n \rfloor}}, 1, \dots, \overbrace{1}^{2^1}, \overbrace{1}^{2^0}}_{\text{only 1's}}$$

The Bitmask Algorithm



The Bitmask Algorithm

- How can we compute $2^{\lfloor \log_2 n \rfloor + 1}$?



The Bitmask Algorithm

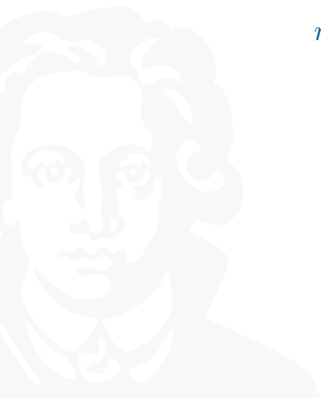
- How can we compute $2^{\lfloor \log_2 n \rfloor + 1}$?
- Count the number ℓ of leading 0's!



The Bitmask Algorithm

- How can we compute $2^{\lfloor \log_2 n \rfloor + 1}$?
- Count the number ℓ of leading 0's!

$$\begin{array}{c}
 n \xrightarrow{\text{binary}} \underbrace{\overbrace{0, \dots, 0}^{2^{W-1}}}_{\ell \text{ 0's}}, \overbrace{1}^{2^{\lfloor \log_2 n \rfloor}}, \underbrace{0/1, \dots, 0/1}_{\text{series of 0's and 1's}}^{\overbrace{2^1} \quad \overbrace{2^0}}
 \end{array}$$



- $\lfloor \log_2 n \rfloor = W - \ell - 1$

- $\lceil \log_2 n \rceil = W - \ell - 1$

$$n \xrightarrow{\text{binary}} \underbrace{0, \dots, 0}_{\ell \text{ 0's}}, \underbrace{1}_{2^{\lfloor \log_2 n \rfloor}}, \underbrace{0/1, \dots, 0/1}_{\text{series of 0's and 1's}}, \underbrace{0/1}_{2^1}, \underbrace{0/1}_{2^0}$$

The Bitmask Algorithm

- How can we compute $2^{\lfloor \log_2 n \rfloor + 1}$?
- Count the number ℓ of leading 0's!

$$\begin{array}{c}
 n \xrightarrow{\text{binary}} \underbrace{\overbrace{0, \dots, 0}^{2^{W-1}}, \overbrace{1}^{2^{\lfloor \log_2 n \rfloor}}, \underbrace{0/1, \dots, 0/1}_{\text{series of 0's and 1's}}}_{\ell \text{ 0's}} \quad \begin{array}{c} \overbrace{0/1}^{2^1} \\ \overbrace{0/1}^{2^0} \end{array}
 \end{array}$$

$$\lfloor \log_2 n \rfloor = W - \ell - 1 \quad \longrightarrow \quad 2^{\lfloor \log_2 n \rfloor + 1} = 1 \ll (W - \ell)$$

The Bitmask Algorithm

- How can we compute $2^{\lfloor \log_2 n \rfloor + 1}$?
- Count the number ℓ of leading 0's!

$$\begin{array}{c}
 n \xrightarrow{\text{binary}} \underbrace{\overbrace{0, \dots, 0}^{2^{W-1}}}_{\ell \text{ 0's}}, \overbrace{1}^{2^{\lfloor \log_2 n \rfloor}}, \underbrace{0/1, \dots, 0/1}_{\text{series of 0's and 1's}}^{\overbrace{2^1} \quad \overbrace{2^0}}
 \end{array}$$

- $\lfloor \log_2 n \rfloor = W - \ell - 1 \quad \longrightarrow \quad 2^{\lfloor \log_2 n \rfloor + 1} = 1 \ll (W - \ell)$
- Algorithm:

The Bitmask Algorithm

- How can we compute $2^{\lfloor \log_2 n \rfloor + 1}$?
- Count the number ℓ of leading 0's!

$$\begin{array}{c}
 n \xrightarrow{\text{binary}} \underbrace{\overbrace{0, \dots, 0}^{2^{W-1}}}_{\ell \text{ 0's}}, \overbrace{1}^{2^{\lfloor \log_2 n \rfloor}}, \underbrace{0/1, \dots, 0/1}_{\text{series of 0's and 1's}}^{\overbrace{2^1} \quad \overbrace{2^0}}
 \end{array}$$

- $\lfloor \log_2 n \rfloor = W - \ell - 1 \quad \longrightarrow \quad 2^{\lfloor \log_2 n \rfloor + 1} = 1 \ll (W - \ell)$
- Algorithm:
 - (1) Compute ℓ and $\mathcal{M} = 2^{\lfloor \log_2 n \rfloor + 1} - 1$

The Bitmask Algorithm

- How can we compute $2^{\lfloor \log_2 n \rfloor + 1}$?
- Count the number ℓ of leading 0's!

$$\begin{array}{c}
 n \xrightarrow{\text{binary}} \underbrace{\overbrace{0, \dots, 0}^{2^{W-1}}}_{\ell \text{ 0's}}, \overbrace{1}^{2^{\lfloor \log_2 n \rfloor}}, \underbrace{0/1, \dots, 0/1}_{\text{series of 0's and 1's}}^{\overbrace{2^1} \quad \overbrace{2^0}}
 \end{array}$$

- $\lfloor \log_2 n \rfloor = W - \ell - 1 \quad \longrightarrow \quad 2^{\lfloor \log_2 n \rfloor + 1} = 1 \ll (W - \ell)$
- Algorithm:
 - (1) Compute ℓ and $\mathcal{M} = 2^{\lfloor \log_2 n \rfloor + 1} - 1$
 - (2) Draw $x \in [0, 2^W)$ and compute $b = x \ \& \ \mathcal{M}$

The Bitmask Algorithm

- How can we compute $2^{\lfloor \log_2 n \rfloor + 1}$?
- Count the number ℓ of leading 0's!

$$\begin{array}{c}
 n \xrightarrow{\text{binary}} \underbrace{\overbrace{0, \dots, 0}^{2^{W-1}}}_{\ell \text{ 0's}}, \overbrace{1}^{2^{\lfloor \log_2 n \rfloor}}, \underbrace{0/1, \dots, 0/1}_{\text{series of 0's and 1's}}^{\overbrace{2^1} \quad \overbrace{2^0}}
 \end{array}$$

- $\lfloor \log_2 n \rfloor = W - \ell - 1 \quad \longrightarrow \quad 2^{\lfloor \log_2 n \rfloor + 1} = 1 \ll (W - \ell)$
- Algorithm:
 - (1) Compute ℓ and $\mathcal{M} = 2^{\lfloor \log_2 n \rfloor + 1} - 1$
 - (2) Draw $x \in [0, 2^W)$ and compute $b = x \ \& \ \mathcal{M}$
 - (3) Return b if $b < n$ else goto (2)

The Bitmask Algorithm

Algorithm:

- (1) Compute ℓ and $\mathcal{M} = 2^{\lfloor \log_2 n \rfloor + 1} - 1$
- (2) Draw $x \in [0, 2^W)$ and compute $b = x \ \& \ \mathcal{M}$
- (3) Return b if $b < n$ else goto (2)



The Bitmask Algorithm

Algorithm:

- (1) Compute ℓ and $\mathcal{M} = 2^{\lceil \log_2 n \rceil + 1} - 1$
- (2) Draw $x \in [0, 2^W)$ and compute $b = x \ \& \ \mathcal{M}$
- (3) Return b if $b < n$ else goto (2)

Efficiency

The Bitmask Algorithm

Algorithm:

- (1) Compute ℓ and $\mathcal{M} = 2^{\lfloor \log_2 n \rfloor + 1} - 1$
- (2) Draw $x \in [0, 2^W)$ and compute $b = x \ \& \ \mathcal{M}$
- (3) Return b if $b < n$ else goto (2)

Efficiency

- b at most $\mathcal{M} = 2^{\lfloor \log_2 n \rfloor + 1} - 1 < 2n$

The Bitmask Algorithm

Algorithm:

- (1) Compute ℓ and $\mathcal{M} = 2^{\lfloor \log_2 n \rfloor + 1} - 1$
- (2) Draw $x \in [0, 2^W)$ and compute $b = x \ \& \ \mathcal{M}$
- (3) Return b if $b < n$ else goto (2)

Efficiency

■ b at most $\mathcal{M} = 2^{\lfloor \log_2 n \rfloor + 1} - 1 < 2n \quad \longrightarrow \quad$ success probability at least $\approx \frac{1}{2}$

The Bitmask Algorithm

Algorithm:

- (1) Compute ℓ and $\mathcal{M} = 2^{\lfloor \log_2 n \rfloor + 1} - 1$
- (2) Draw $x \in [0, 2^W)$ and compute $b = x \ \& \ \mathcal{M}$
- (3) Return b if $b < n$ else goto (2)

Efficiency

- b at most $\mathcal{M} = 2^{\lfloor \log_2 n \rfloor + 1} - 1 < 2n \quad \longrightarrow \quad$ success probability at least $\approx \frac{1}{2}$
- At most 2 rounds in expectation

The Bitmask Algorithm

Algorithm:

- (1) Compute ℓ and $\mathcal{M} = 2^{\lfloor \log_2 n \rfloor + 1} - 1$
- (2) Draw $x \in [0, 2^W)$ and compute $b = x \ \& \ \mathcal{M}$
- (3) Return b if $b < n$ else goto (2)

Efficiency

- b at most $\mathcal{M} = 2^{\lfloor \log_2 n \rfloor + 1} - 1 < 2n \quad \longrightarrow \quad$ success probability at least $\approx \frac{1}{2}$
- At most 2 rounds in expectation
- No integer division at all

The Bitmask Algorithm

Algorithm:

- (1) Compute ℓ and $\mathcal{M} = 2^{\lfloor \log_2 n \rfloor + 1} - 1$
- (2) Draw $x \in [0, 2^W)$ and compute $b = x \ \& \ \mathcal{M}$
- (3) Return b if $b < n$ else goto (2)

Efficiency

- b at most $\mathcal{M} = 2^{\lfloor \log_2 n \rfloor + 1} - 1 < 2n \rightarrow$ success probability at least $\approx \frac{1}{2}$
- At most 2 rounds in expectation
- No integer division at all
- Computation of leading 0's requires `clz` instruction/algorithm

The Bitmask Algorithm

Algorithm:

- (1) Compute ℓ and $\mathcal{M} = 2^{\lfloor \log_2 n \rfloor + 1} - 1$
- (2) Draw $x \in [0, 2^W)$ and compute $b = x \ \& \ \mathcal{M}$
- (3) Return b if $b < n$ else goto (2)

Efficiency

- b at most $\mathcal{M} = 2^{\lfloor \log_2 n \rfloor + 1} - 1 < 2n \rightarrow$ success probability at least $\approx \frac{1}{2}$
- At most 2 rounds in expectation
- No integer division at all
- Computation of leading 0's requires `clz` instruction/algorithm
- Roughly as expensive as a `div` instruction

Lemire's Algorithm

Multiply-And-Shift



Multiply-And-Shift

$$(\text{rand}() \cdot n) \gg W$$



Multiply-And-Shift

$$(\text{rand}() \cdot n) \gg W = (\text{rand}() \cdot n) \div 2^W$$



Lemire's Algorithm

The Algorithm



Lemire's Algorithm

The Algorithm



4 Conclusion



Conclusion

Summary



expected number of integer division operations	maximum number of integer division operations	Unbiased?
--	---	-----------



	expected number of integer division operations	maximum number of integer division operations	Unbiased?
Modulo Reduction	1	1	X



	expected number of integer division operations	maximum number of integer division operations	Unbiased?
Modulo Reduction	1	1	X
Multiply-and-Shift	0	0	X

	expected number of integer division operations	maximum number of integer division operations	Unbiased?
Modulo Reduction	1	1	X
Multiply-and-Shift	0	0	X
OpenBSD	2	2	✓

Conclusion

Summary

	expected number of integer division operations	maximum number of integer division operations	Unbiased?
Modulo Reduction	1	1	✗
Multiply-and-Shift	0	0	✗
OpenBSD	2	2	✓
Java	$\frac{2^W}{2^W - (2^W \bmod n)}$	∞	✓

	expected number of integer division operations	maximum number of integer division operations	Unbiased?
Modulo Reduction	1	1	✗
Multiply-and-Shift	0	0	✗
OpenBSD	2	2	✓
Java	$\frac{2^W}{2^W - (2^W \bmod n)}$	∞	✓
Bitmask	0	0	✓

Conclusion

Summary

	expected number of integer division operations	maximum number of integer division operations	Unbiased?
Modulo Reduction	1	1	✗
Multiply-and-Shift	0	0	✗
OpenBSD	2	2	✓
Java	$\frac{2^W}{2^W - (2^W \bmod n)}$	∞	✓
Bitmask	0	0	✓
Lemire	$\frac{n}{2^W}$	1	✓

End of Talk

